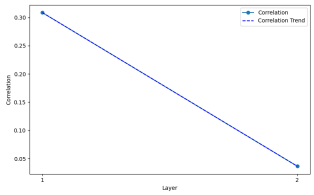

(a) Layers: 2, Heads: 8, Attn-only: False, Adam

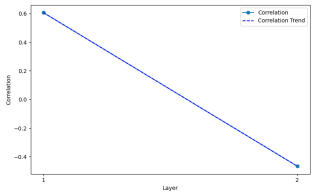

(b) Layers: 2, Heads: 16, Attn-only: False, Adam

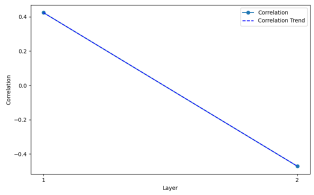

(c) Layers: 2, Heads: 8, Attn-only: False, AdamW

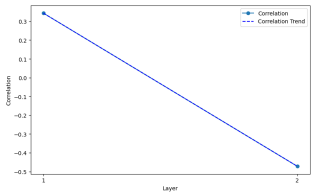

(d) Layers: 2, Heads: 16, Attn-only: False, AdamW

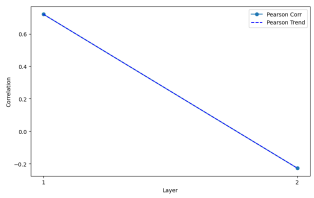

(a) Layers: 2, Heads: 8, Attn-only: True, Adam

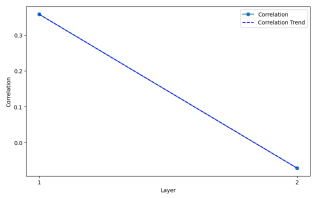

(b) Layers: 2, Heads: 16, Attn-only: True, Adam

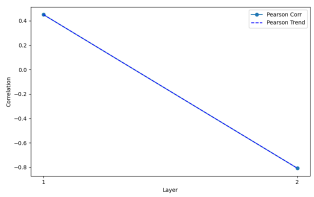

(c) Layers: 2, Heads: 8, Attn-only: True, AdamW

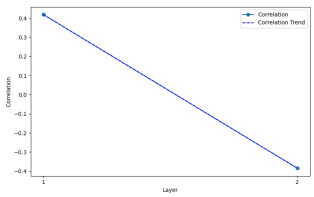

(d) Layers: 2, Heads: 16, Attn-only: True, AdamW

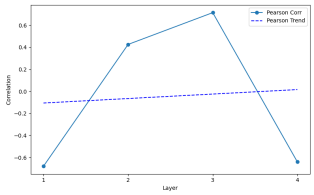

(a) Layers: 4, Heads: 8, Attn-only: False, Adam

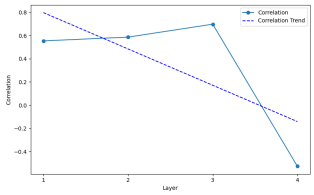

(b) Layers: 4, Heads: 16, Attn-only: False, Adam

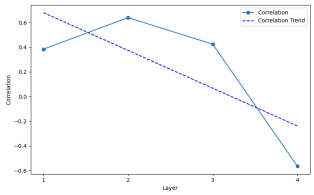

(c) Layers: 4, Heads: 8, Attn-only: False, AdamW

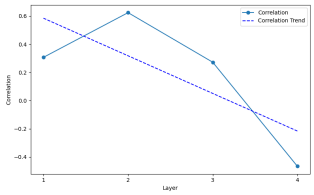

(d) Layers: 4, Heads: 16, Attn-only: False, AdamW

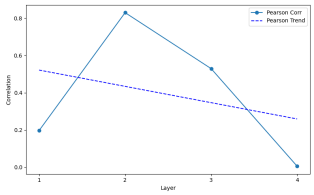

(a) Layers: 4, Heads: 8, Attn-only: True, Adam

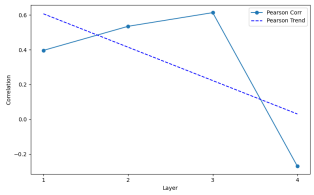

(b) Layers: 4, Heads: 16, Attn-only: True, Adam

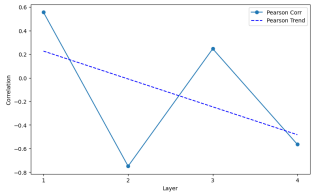

(c) Layers: 4, Heads: 8, Attn-only: True, AdamW

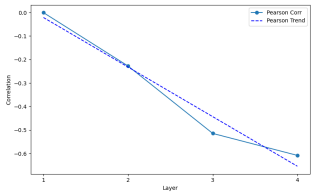

(d) Layers: 4, Heads: 16, Attn-only: True, AdamW

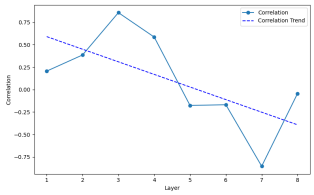

(a) Layers: 8, Heads: 8, Attn-only: False, Adam

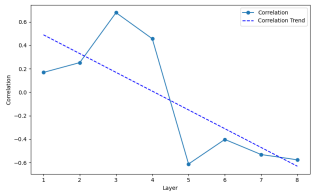

(b) Layers: 8, Heads: 16, Attn-only: False, Adam

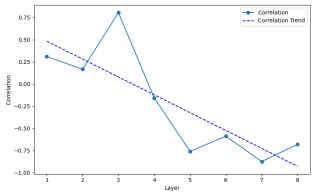

(c) Layers: 8, Heads: 8, Attn-only: False, AdamW

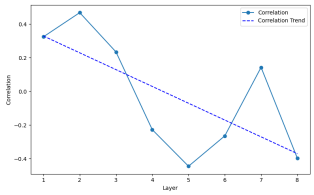

(d) Layers: 8, Heads: 16, Attn-only: False, AdamW

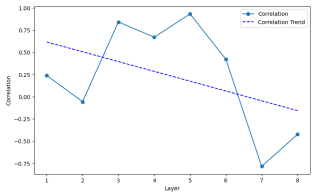

(a) Layers: 8, Heads: 8, Attn-only: True, Adam

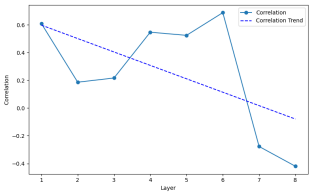

(b) Layers: 8, Heads: 16, Attn-only: True, Adam

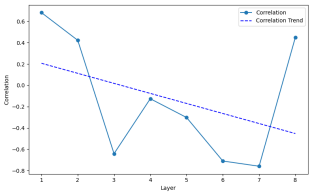

(c) Layers: 8, Heads: 8, Attn-only: True, AdamW

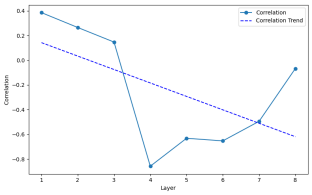

(d) Layers: 8, Heads: 16, Attn-only: True, AdamW

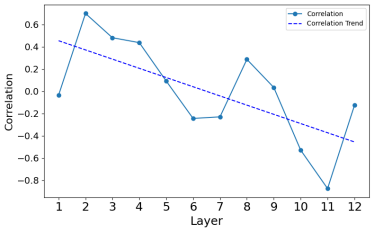

(a) Layers: 12, Heads: 12, Attn-only: False, Adam

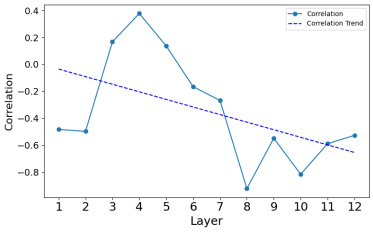

(b) Layers: 12, Heads: 12, Attn-only: False, AdamW

Listing 1. TransformerLens HookedTransformerConfig template for MLP architectures

```
{
'act_fn': 'gelu',
 'attention_dir': 'causal',
 'attn_only': False,
 'attn_scale': np.float64(8.0),
 'attn_scores_soft_cap': -1.0,
 'attn_types': None,
 'checkpoint_index': None,
 'checkpoint_label_type': None,
 'checkpoint_value': None,
 'd_head': 64,
 'd_mlp': 2048,
 'd_model': 512,
 'd_vocab': 48262,
 'd_vocab_out': 48262,
 'decoder_start_token_id': None,
 'default_prepend_bos': True,
 'device': device(type='cuda'),
 'dtype': torch.float32,
 'eps': 1e-05,
 'experts_per_token': None,
 'final_rms': False,
 'from_checkpoint': False,
 'gated_mlp': False,
 'init_mode': 'gpt2',
 'init_weights': True,
 'initializer_range': np.float64(0.035355339059327376),
 'load_in_4bit': False,
 'model_name': 'GELU_8L512W_C4_Code_8H',
 'n_ctx': 1024,
 'n_devices': 1,
 'n_heads': [NUM\_HEADS],
 'n_key_value_heads': None,
 'n_layers': [NUM\_LAYERS],
 'n_params': 25165824,
 'normalization_type': 'LN',
 'num_experts': None,
 'original_architecture': None,
 'output_logits_soft_cap': -1.0,
 'parallel_attn_mlp': False,
 'positional_embedding_type': 'standard',
 'post_embedding_ln': False,
 'relative_attention_max_distance': None,
 'relative_attention_num_buckets': None,
 'rotary_adjacent_pairs': False,
 'rotary_base': 10000,
 'rotary_dim': None,
 'scale_attn_by_inverse_layer_idx': False,
 'seed': 2,
 'tie_word_embeddings': False,
 'tokenizer_name': 'NeelNanda/gpt-neox-tokenizer-digits',
 'tokenizer_prepends_bos': False,
 'trust_remote_code': False,
 'use_attn_in': False,
 'use_attn_result': False,
 'use_attn_scale': True,
 'use_hook_mlp_in': False,
 'use_hook_tokens': False,
 'use_local_attn': False,
 'use_normalization_before_and_after': False,
 'use_split_qkv_input': False,
 'window_size': None
 }
```